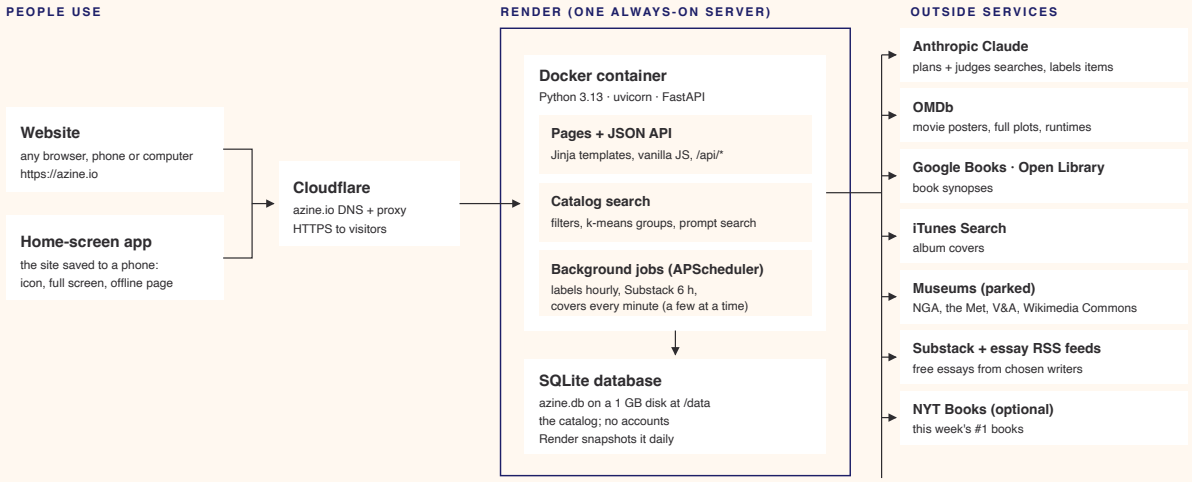
AZINE

System summary

How azine works, on one page

OCTOBER 29, 2026
AZINE.IO · MAIN @ 66CEE45

azine is a catalog of essays, movies, albums and books (E-M-A-B), the same for everyone: no accounts, no personalized picks, no voting. No feeds, no influencers, no AI-written content: every entry is a real work from a curated catalog. Find things by category, by when the work was made, and by saying what you want.



How it works

1. **Category, date and country**: any of the four, the years the work was made, and (not essays) where it's from. A plain query, no Claude.
2. **Groups**: each category clustered (k-means over TF-IDF embeddings of genres, subjects, tones, synopsis) into named groups of similar work. No Claude.
3. **Prompt search**, inside those filters: Claude Sonnet plans it (the distinct things asked for, as labels, keywords, typical works, groups); the server lines up every work in the filters, each part of the prompt getting its share (after PinnerSage); Haiku scores each 0–100 and every 75+ comes back.
4. **Background jobs**: label new items hourly, Substack every 6 h, covers a few at a time. Parked: museum art (old board at tag `personal-recommender-final`).

Running it

Catalog: ~1,900 movies, ~2,000 books, ~1,800 albums (incl. ~1,300 Latin, Korean, African), ~2,000 essays (+ live Substack/RSS).

Quotas: OMDb 1,000/day; Google Books needs its own free key.

Memory: ~310 MB of 512 MB with the old board (re-measure); move to 1 CPU / 2 GB if logs show it running out.

Backups: daily disk snapshots on Render, kept ≥7 days.

Deploy: push to `main` (~1 min downtime). **More:** `DEPLOY.md`, `README.md`, and the full architecture PDF.

Stack

PART	WHAT
Server	Python 3.13, FastAPI, uvicorn, SQLAlchemy, Alembic, APScheduler, scikit-learn
Data	SQLite file on a persistent disk; migrations run on start
Front end	E M A B down the left; a click grows the letter into a title card, then its catalog (5 columns, a timeline, "Vibes" search, Met art while it runs). One template + vanilla JS + one stylesheet
Hosting	Render (Docker, 0.5 CPU / 512 MB, 1 GB disk), deploys on every push to `main`
Domain	azine.io on Cloudflare DNS → Render; HTTPS everywhere
Phones	Home-screen app: the site saved to a phone's home screen from the browser, with an icon, full screen and an offline page. No App Store app
AI	Claude plans (Sonnet) and judges (Haiku) searches and labels items. It never writes what you read.

Key decisions

- **One server + SQLite**: simplest thing that works at this size; scales by upgrading the instance.
- **Render** over Fly.io: a dashboard instead of a CLI; Free tier ruled out (sleeps, no disk).
- **A catalog, not a recommender**: no sign-up between a visitor and the work.
- **A 75% bar, not a top 10**: Claude judges fit, because similarity can't tell a mood from a subject.
- **TF-IDF, not a neural model**: separates genres well at this catalog size without a multi-GB dependency.