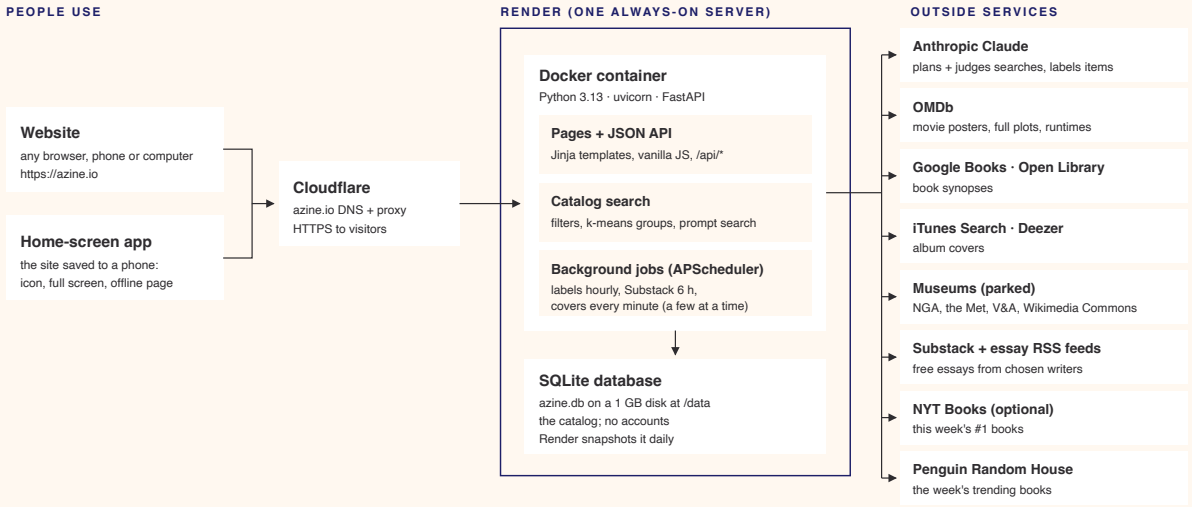
AZINE

System summary

How azine works, on one page

NOVEMBER 9, 2026
AZINE.IO · MAIN @ 98552E2

azine is a catalog of essays, movies, albums and books (E-M-A-B), the same for everyone: no accounts, no personal picks, nothing AI-written. Browse by category, year and country, or describe what you want.



How it works

1. **Filters**: categories, years made, countries (not essays). A plain query, no Claude.
2. **Groups**: each category clustered (k-means over TF-IDF of labels and synopsis) into named groups. No Claude.
3. **Search**: Sonnet plans the prompt; the server lines up every work in the filters by similarity; Haiku scores each 0–100 and every 75+ comes back.
4. **Background jobs**: labels hourly, Substack every 6 h, trending books weekly, covers a few at a time.

Running it

Catalog: ~1,900 movies, ~3,500 books, ~1,750 albums, ~1,900 essays (+ live feeds).

Limits: OMDb 1,000/day; 512 MB memory (~310 MB used, re-measure).

Backups: daily Render disk snapshots, kept ≥7 days.

Deploy: push to `main` (~1 min downtime). See `DEPLOY.md`, `README.md`.

Stack

PART	WHAT
Server	Python 3.13, FastAPI, SQLAlchemy, Alembic, APScheduler, scikit-learn
Data	SQLite on a persistent disk; migrations on start
Front end	One template, vanilla JS, one stylesheet: E M A B, a title card, then the catalog and "Vibes" search
Hosting	Render (Docker, 0.5 CPU / 512 MB, 1 GB disk); azine.io via Cloudflare
Phones	Home-screen web app; no App Store app
AI	Claude plans and judges searches and labels items. It never writes what you read

Key decisions

- **One server + SQLite**: simplest at this size; scale by upgrading.
- **Render**: a dashboard and GitHub deploys; the free tier sleeps and has no disk.
- **A catalog, not a recommender**: no sign-up before the work.
- **A 75% bar, not a top 10**: similarity can't tell a mood from a subject; Claude can.
- **TF-IDF**: good enough at this size, no multi-GB model.